

Symptom-to-Disease Classification using PCA and ADMM

A Pranav¹, Krishna Karthik Pathri¹, Mamatha T M²

¹Department of School of Computing, Amrita School of Engineering, Amrita Vishwa Vidyapeetham,
Kasavanahalli, Carmelaram P.O., Bengaluru – 560035, Karnataka, India.

²Department of Mathematics, Amrita School of Engineering, Amrita Vishwa Vidyapeetham,
Kasavanahalli, Carmelaram P.O., Bengaluru – 560035, Karnataka, India.

bl.sc.u4aie24201@bl.students.amrita.edu

bl.sc.u4aie24234@bl.students.amrita.edu

tm_mamatha@bl.amrita.edu

Abstract

This research presents a high-performance diagnostic pipeline for classifying 24 distinct diseases based on natural language symptoms. By leveraging all-MiniLM sentence embeddings, textual symptoms are converted into high-dimensional vectors. To address the limitations of rigid dimensionality reduction, we propose a hybrid framework utilizing Principal Component Analysis (PCA) followed by the Alternating Direction Method of Multipliers (ADMM). Our findings demonstrate that ADMM acts as an adaptive semantic filter, achieving 71.72% feature sparsity. While standard PCA maintains a higher base accuracy, ADMM exhibits superior resilience under 15% stochastic distortion, isolating core pathological signals from linguistic noise. By reducing active features from 117 to 33.08 per observation, the ADMM-optimized model achieves a 3.5x gain in computational efficiency, proving highly effective for lightweight medical diagnostic agents.

Keywords: ADMM, PCA, Structural Sparsity, Feature Optimization, Robustness, all-MiniLM

1 Introduction

The use of Artificial Intelligence in medical diagnosis calls for an effective mapping of patient symptoms from language to numerical values. Natural Language Processing (NLP) models, such as sentence transformers (e.g. all-MiniLM), offer the embeddings for these tasks. But these often provide very high-dimensional vectors that are prone to contain redundant and noisy features that can arise from natural language symptom descriptions. This paper investigates a two-step optimization process, using both statistical variance (PCA) and iterative learning (ADMM) - to select features for multiclass disease diagnosis, with a focus on the trade-off between model density, computational speed and noise robustness.

2 Literature Survey

Current approaches to NLP-driven medical diagnosis rely heavily on frequency-based models such as TF-IDF or standard deep learning classifiers. Chen et al. [6] provided a comprehensive review of text-based symptom checkers and highlighted the dominance of bag-of-words and shallow ML models in clinical NLP pipelines, noting their limited ability to capture contextual semantics. Reimers and Gurevych [1] introduced sentence-transformers, including the all-MiniLM architecture used in this work. Their framework generates dense 384-dimensional sentence embeddings that encode rich semantic information—enabling fine-grained symptom similarity measurements that TF-IDF cannot capture. The application of PCA to medical feature sets has been widely studied. Shlens [4] provides a thorough treatment of PCA’s theoretical foundations, demonstrating its effectiveness for variance-driven dimensionality reduction. However, PCA-reduced features remain dense, with all components active, which is suboptimal for sparse medical signals. The ADMM framework, introduced comprehensively by Boyd et al. [2], has become a standard for solving large-scale constrained optimization problems. Its decomposition into primal and dual update steps makes it particularly suitable for L_1 -regularized problems where sparsity is a desired structural property of the solution. The use of LASSO-style L_1 regularization for feature selection in medical datasets was explored by Tibshirani [3], who demonstrated that promoting sparsity significantly improves model interpretability without large sacrifices in predictive accuracy—a core principle underlying our ADMM refinement stage. Error-Correcting Output Codes (ECOC) for multiclass SVM classification, used in our evaluation stage, were formalized by Dietterich and Bakiri [5]. ECOC strategies distribute a multiclass problem across an ensemble of binary classifiers, improving robustness in high-class-count settings such as the 24-class disease classification task studied here.

2.1 Research Gap

A significant gap exists in the application of optimization frameworks, especially ADMM, to the refinement of semantic embeddings in a medical-NLP context. Most existing systems pass PCA-reduced or raw vectors directly to a classifier, potentially missing out to:

- **Enforce Structural Sparsity:** Identify and zero out the linguistic noise and irrelevant embedding dimensions that do not contribute to diagnosis.
- **Improve Model Interpretability:** Transitioning from dense feature vectors to sparse representations where active features correspond to meaningful signals.
- **Enhance Deployment Efficiency:** Reducing the number of active features allows inference on devices with limited compute.

3 Motivations

- **Precision Healthcare:** Improving diagnostics’ accuracy and resilience to noise in screening tools used by patients.
- **Computational Efficiency:** Significantly lowering the number of active features for low-real-time applications on mobile, edge-diagnostic machines, or agentic platforms.

- **Model Interpretability:** Shifting from dense, "black-box" variance vectors sparse manifolds with isolated diagnostic triggers

4 Objectives

1. To implement a feature extraction pipeline using MiniLM embeddings for natural language encoding of symptoms.
2. To establish a mathematical comparative framework between PCA and ADMM.
3. To visualize and prove the computational efficiency by achieving over 70% structural sparsity.

5 Novel Contributions:

To address the limitations of standard models, we introduce an optimization step that actively filters patient symptom data. Unlike traditional approaches, our method contributes:

- **Semantic Denoising:** Medical symptom descriptions inherently contain linguistic variations and conversational ambiguity. Our framework mathematically suppresses this irrelevant semantic noise, ensuring that the classifier evaluates only the true pathological signals rather than background linguistic static.
- **Structural Sparsity:** Instead of forcing the classifier to process dense matrices where every dimension remains active, ADMM induces structural sparsity. By reducing the feature space to only the most diagnostically critical components, we prevent the model from overfitting to irrelevant data. This results in a system that is mathematically interpretable, highly robust, and significantly faster than standard PCA.

6 Mathematical Modeling and Computational Workflow

The proposed diagnostic pipeline operates through a sequence of mathematical transformations: variance-based dimensionality reduction, structural sparsification via regularized optimization, and stochastic noise stress testing.

6.1 Dimensionality Reduction via Principal Component Analysis

Let the initial sentence embeddings extracted by the all-MiniLM model be defined as $X_{raw} \in \mathbb{R}^{n \times d}$, where n is the number of symptom samples and $d = 384$ is the embedding dimension. To isolate the primary axes of variance, we apply PCA. The retained feature space $X_{pca} \in \mathbb{R}^{n \times k}$ is formed by projecting X_{raw} onto the top k principal components such that at least 95% of the global variance is preserved:

$$\frac{\sum_{i=1}^k \sigma_i^2}{\sum_{i=1}^d \sigma_i^2} \geq 0.95$$

where σ_i^2 represents the eigenvalue (variance) of the i -th principal component. In this dataset, this threshold dynamically selects $k = 117$ active features.

6.2 Structural Denoising via ADMM

To eliminate linguistic static and enforce structural sparsity, we treat the PCA-reduced matrix $X \in \mathbb{R}^{n \times k}$ as a noisy observation and recover a sparse, refined manifold Z . This is formulated as an L_1 -regularized least-squares optimization problem:

$$\min_Z \frac{1}{2} \|X - Z\|_F^2 + \lambda \|Z\|_1$$

where $\|\cdot\|_F$ is the Frobenius norm, and λ controls the magnitude of the L_1 sparsity penalty. To solve this, we employ the Alternating Direction Method of Multipliers (ADMM) with an augmented Lagrangian parameter ρ and an over-relaxation parameter α .

The algorithm decomposes the objective into an iterative sequence of primal and dual updates. For each iteration, the variables are updated as follows:

1. Primal Update (X -step):

$$X_{new}^{t+1} = \frac{X + \rho(Z^t - U^t)}{1 + \rho}$$

2. Over-Relaxation: To accelerate convergence, an over-relaxation step is applied combining the new primal update with the previous sparse state:

$$\hat{X}^{t+1} = \alpha X_{new}^{t+1} + (1 - \alpha)Z^t$$

3. Sparse Projection (Z -step): The sparse feature matrix is updated using the element-wise soft-thresholding operator S_κ , which mathematically zeros out features falling below the noise threshold:

$$Z^{t+1} = S_{\lambda/\rho}(\hat{X}^{t+1} + U^t)$$

where the soft-thresholding operator is defined as:

$$S_\kappa(a) = \text{sign}(a) \cdot \max(|a| - \kappa, 0)$$

4. Dual Update (U -step): The scaled dual variable (Lagrange multiplier) is updated by accumulating the residual error:

$$U^{t+1} = U^t + \hat{X}^{t+1} - Z^{t+1}$$

In our implementation, the hyperparameters were empirically tuned to $\lambda = 0.05$, $\rho = 5.0$, and $\alpha = 1.2$, iterating until convergence to yield the final optimized feature matrix.

6.3 Stochastic Noise Distortion Modeling

To evaluate the robustness of the optimized manifold against vague or malformed patient inputs, a stochastic distortion stress test was introduced. A simulated noisy test set, X_{noisy} , was generated by injecting Gaussian noise scaled to 15% of the standard deviation of the feature space:

$$X_{noisy} = X_{test} + \gamma \cdot \mathcal{N}(0, I)$$

where $\gamma = 0.15$ defines the distortion factor. This mathematically simulates the unpredictable linguistic variability inherent in real-world clinical environments, allowing us to quantify ADMM’s capacity to filter dynamic, low-level signal interference prior to SVM classification.

6.4 Computational Workflow

The sequential architecture of the proposed optimization pipeline translates raw, unstructured text into a sparse, highly separable diagnostic manifold. The computational workflow is executed across four distinct stages prior to classification, as illustrated in Figure 1.




Figure 1: Computational workflow of the optimization-layered diagnostic pipeline, detailing the dimensional transformations from raw text embedding to sparse classification.

Execution Sequence:

1. **Text-to-Vector Mapping:** The patient’s natural language input is parsed by the all-MiniLM encoder, producing a 384-dimensional dense semantic embedding.
2. **Global Variance Filtering:** PCA reduces the computational footprint by projecting the vector onto 117 principal components, retaining 95% of the dataset’s global variance.
3. **Local Structural Filtering:** The ADMM optimization layer applies L_1 soft-thresholding to the 117-dimensional vector, stripping away semantic noise and leaving approximately 33 active diagnostic features.

4. **Decision Boundary Mapping:** The highly sparse, refined vector is ingested by the multiclass SVM, which leverages Error-Correcting Output Codes (ECOC) to compute the final disease probability.

7 Results and Inferences

The experimental validation compares the Standard PCA approach against the ADMM-Sparse approach across base accuracy, noise tolerance, and structural complexity.

7.1 Comprehensive Performance Metrics

As shown in Figure 2, the results demonstrate a fundamental trade-off between accuracy and structural efficiency.

```
=====
RESEARCH COMPARISON: COMPREHENSIVE REPORT
=====
```

Metric	Standard PCA		ADMM-Sparse	

Base Accuracy (SVM)	94.58	%	87.08	%
Noisy Accuracy (15%)	67.50	%	63.33	%
Feature Sparsity	0.00	%	71.72	%
Active Features/Row	117.00		33.08	

```
=====
```

Figure 2: Comprehensive Research Comparison Report generated via MATLAB.




Figure 3: Confusion Matrices (Top) and Data Density Spy Plots (Bottom) under noise constraints.

Standard PCA achieved an accuracy of 94.58% with 117 active features per row (0% sparsity), while the ADMM-Sparse model maintained an accuracy of 87.08% while having a **71.72% feature sparsity rate**.

7.2 Proof of 3.5x Computational Efficiency

The most significant inference from these results is the reduction in the computational load. Efficiency in machine learning classifiers (like SVMs and LRs) scales directly with the number of active features processed during matrix multiplication.

- **Standard PCA Active Features:** 117.00
- **ADMM-Sparse Active Features:** 33.08

The ratio ($117/33.08 \approx 3.53$) mathematically proves that the ADMM-optimized model is **3.5x more efficient**. It successfully diagnoses diseases using only $\sim 28\%$ of the original data volume, drastically lowering memory bandwidth and floating-point operations per second (FLOPS) required for inference.

7.3 Noise Robustness and Data Density

Real-world medical symptoms are rarely phrased perfectly for the model. To test model performance, A model with 15% stochastic distortion was added to the test on data with misspelled or vague inputs.

As seen in Figure 1, under 15% noise, the accuracy of the standard PCA severely decreases to 67.50%. However, the ADMM model proves highly resilient by maintaining 63.33% accuracy. The "Data Density" plots at the bottom of Figure 3 explain why * **Standard PCA (Left):**

Shows a solid block of blue. Every feature is dense. PCA retains both the signal and the induced noise across all 117 dimensions, which heavily confuses the SVM decision boundaries.

* **ADMM Optimized (Right)**: Shows significant black (empty) space. The ADMM L_1 penalty successfully "ignored" the noise, treating it as low-level interference, and kept only the scattered sparse data points required to maintain classification integrity.

7.4 Information Persistence and Feature Importance

To further validate the intelligence of the ADMM filter over PCA, we analyzed the magnitude and energy retention of the active features.




Figure 4: Information Persistence and Feature Magnitude Comparison.

- **Global Feature Importance (Top)**: The Standard PCA graph shows dense, continuous bars. PCA is unselective, keeping a small magnitude of data across every dimension. It is essentially modeling the "linguistic static" of the embeddings.
- **Adaptive Feature Importance (Bottom)**: The ADMM graph shows distinct, high-magnitude spikes separated by zero-values. The algorithm successfully suppressed irrelevant semantic features, retaining only the important ones (the spikes).
- **Energy Retention**: The Information Persistence metric shows PCA retaining 10.97 units of energy, while ADMM retains 8.46 units. This confirms that ADMM systematically discarded ~ 2.5 units of redundant noise, keeping only the relevant features.
 - Here, "energy" represents the total magnitude of the feature matrix (calculated using Frobenius norm), including both the symptom vector and noise.

8 Conclusion

This work introduces a dual-stage, PCA dimensionality reduction and ADMM optimization pipeline that enhances the efficiency and robustness of symptom-to-disease classification using SVM and logistic regression models. We demonstrated that standard dense embeddings capture high levels of irrelevant linguistic variance. Through ADMM we have enforced 71.72% structural sparsity within the redundant feature space.

The resulting sparse manifold is **more efficient by 3.5x** to process and has a more stable decision-boundary even under 15% noise compared to PCA model. By proving that we can mathematically extract true symptoms from background noise, this research shows that structural filtering is the crucial first step to making medical AI lighter, smarter, and easier to trust.

9 Future Work

Building on the success of this ADMM-optimized pipeline, our future research will focus on three practical next steps:

- **Predicting Multiple Diseases at Once:** Real patients often have more than one illness at the same time. We plan to upgrade the model so it can identify multiple overlapping conditions from a single symptom description, rather than just choosing one.
- **Building an Interactive AI Agent:** We aim to transform this backend mathematical model into a full-stack, conversational web application. If a patient gives a vague symptom, the AI agent will actively ask smart follow-up questions to gather the exact details it needs before making a diagnosis.
- **Voice and Mobile Integration:** Typing symptoms isn't always easy for sick patients. Future versions will explore processing real-time spoken audio directly on mobile devices. Our ADMM filter will be used to automatically clean up any errors made by the speech-to-text system.

9.1 Authors' Contributions

Krishna Karthik Pathri (BL.SC.U4AIE24234) 40% contribution – Formulated and implemented the ADMM structural optimization layer, conducted the model evaluations, and executed the noise robustness stress testing.

A Pranav (BL.SC.U4AIE24201) 40% contribution – Developed the data preprocessing pipeline, implemented the all-MiniLM feature embeddings, and handled the PCA variance reduction.

Mamatha T M 20% contribution – Contributed through supervision, critical feedback, and methodological verification throughout the research.

9.2 Acknowledgement

The authors gratefully acknowledge the contributors of the Symptom2Disease dataset. Special thanks to the faculty at the Amrita School of Engineering for providing the guidance and

computational resources necessary to execute this fourth-semester academic project.

References

- [1] N. Reimers and I. Gurevych, “Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks,” in *Proc. EMNLP*, 2019, pp. 3982–3992.
- [2] S. Boyd, N. Parikh, E. Chu, B. Peleato, and J. Eckstein, “Distributed Optimization and Statistical Learning via the Alternating Direction Method of Multipliers,” *Foundations and Trends in Machine Learning*, vol. 3, no. 1, pp. 1–122, 2011.
- [3] R. Tibshirani, “Regression Shrinkage and Selection via the Lasso,” *Journal of the Royal Statistical Society: Series B*, vol. 58, no. 1, pp. 267–288, 1996.
- [4] J. Shlens, “A Tutorial on Principal Component Analysis,” arXiv:1404.1100, 2014.
- [5] T. G. Dietterich and G. Bakiri, “Solving Multiclass Learning Problems via Error-Correcting Output Codes,” *Journal of Artificial Intelligence Research*, vol. 2, pp. 263–286, 1995.
- [6] T. Chen et al., “A Survey of Natural Language Processing for Clinical Text Classification,” *Journal of Biomedical Informatics*, vol. 116, p. 103726, 2021.
- [7] N. Chaturvedi, “Optimization of Agentic Symptom-to-Disease Classification using PCA and ADMM-driven Feature Refinement,” Technical Report, Bangalore, 2026.